

3. Install Docker

- `sudo dnf update -y && sudo dnf install -y docker && sudo systemctl enable docker && sudo systemctl start docker && sudo docker run hello-world`

```
ec2-user@ip-172-31-41-14:~$ sudo dnf update -y && sudo dnf install -y docker && sudo systemctl enable docker && sudo s
systemctl start docker && sudo docker run hello-world
Amazon Linux 2023 Kernel Livepatch repository                               69 kB/s | 11 kB    00:00
Dependencies resolved.
Nothing to do.
Complete!
Last metadata expiration check: 0:00:01 ago on Wed Sep 18 15:59:52 2024.
Dependencies resolved.
```

Package	Architecture	Version	Repository	Size
Installing:				
docker	x86_64	25.0.6-1.amzn2023.0.2	amazonlinux	44 M
Installing dependencies:				
containerd	x86_64	1.7.20-1.amzn2023.0.1	amazonlinux	35 M
iptables-libs	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	401 k
iptables-nft	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	183 k
libcgroup	x86_64	3.0-1.amzn2023.0.1	amazonlinux	75 k
libnetfilter_conntrack	x86_64	1.0.8-2.amzn2023.0.2	amazonlinux	58 k
libnftnl	x86_64	1.0.1-19.amzn2023.0.2	amazonlinux	30 k
libnftnl	x86_64	1.2.2-2.amzn2023.0.2	amazonlinux	84 k
pigz	x86_64	2.5-1.amzn2023.0.3	amazonlinux	83 k
runc	x86_64	1.1.13-1.amzn2023.0.1	amazonlinux	3.2 M

```
Transaction Summary
=====
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
```

```
ec2-user@ip-172-31-41-14:~$ Complete!
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
c1ec31eb5944: Pull complete
Digest: sha256:91fb4b041da273d5a3273b6d587d62d518300a6ad268b28628f74997b93171b2
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/

[ec2-user@ip-172-31-41-14 ~]$
```

Then, configure cgroup in a daemon.json file. This allows kubernetes to manage host more efficiently.

- `cd /etc/docker`
- `cat <<EOF | sudo tee /etc/docker/daemon.json`
`{`
 `"exec-opts": ["native.cgroupdriver=systemd"]`
`}` EOF

```
[ec2-user@ip-172-31-41-14 docker]$ sudo tee /etc/docker/daemon.json <<EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
[ec2-user@ip-172-31-41-14 docker]$ |
```

- `sudo systemctl daemon-reload`
- `sudo systemctl restart docker`

```
[ec2-user@ip-172-31-41-14 docker]$ sudo systemctl daemon-reload
sudo systemctl restart docker
[ec2-user@ip-172-31-41-14 docker]$ |
```

4. Install Kubernetes

Note: I'm directly installing binary package you may install from package repository of your distribution

Install CNI plugins (required for most pod network):

```
CNI_PLUGINS_VERSION="v1.3.0" ARCH="amd64"
DEST="/opt/cni/bin"
sudo mkdir -p "$DEST"
curl -L
"https://github.com/containernetworking/plugins/releases/download/${CNI_PLUGINS_
VERSION}/cni-plugins-linux-${ARCH}-${CNI_PLUGINS_VERSION}.tgz" | sudo tar -C
"$DEST" -xz
```

```
[ec2-user@ip-172-31-41-14 docker]$ CNI_PLUGINS_VERSION="v1.3.0" ARCH="amd64"
DEST="/opt/cni/bin"
sudo mkdir -p "$DEST"
curl -L "https://github.com/containernetworking/plugins/releases/download/${CNI_PLUGINS_VERSION}/cni-plugins-linux-${ARCH}-${CNI_PLUGINS_VERSION}.tgz" | sudo tar -C "$DEST" -xz
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
  0     0    0     0    0     0      0  0 --:--:-- --:--:-- --:--:--     0
100 43.2M 100 43.2M    0     0 42.1M      0  0:00:01  0:00:01 --:--:-- 52.2M
[ec2-user@ip-172-31-41-14 docker]$
```

Define the directory to download command files:

```
DOWNLOAD_DIR="/usr/local/bin"
sudo mkdir -p "$DOWNLOAD_DIR"
```

```
[ec2-user@ip-172-31-41-14 docker]$ DOWNLOAD_DIR="/usr/local/bin"
sudo mkdir -p "$DOWNLOAD_DIR"
```

Optionally install crictl (required for interaction with the Container Runtime Interface (CRI), optional for kubeadm):

```
CRICTL_VERSION="v1.31.0"
```

```
ARCH="amd64"
```

```
curl -L
```

```
"https://github.com/kubernetes-sigs/cri-tools/releases/download/${CRICTL_VERSION}/
```

```
c rictl-${CRICTL_VERSION}-linux-${ARCH}.tar.gz" | sudo tar -C $DOWNLOAD_DIR -xz
```

```
[ec2-user@ip-172-31-41-14 docker]$ CRICTL_VERSION="v1.31.0" ARCH="amd64"
curl -L "https://github.com/kubernetes-sigs/cri-tools/releases/download/${CRICTL_VERSION}/c rictl-${CRICTL_VERSION}-linux-${ARCH}.tar.gz" | sudo tar -C "$DOWNLOAD_DIR" -xz
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
0	0	0	0	0	--:--:--	--:--:--	0
100	17.5M	100	17.5M	0	0	31.3M	66.4M

Install kubeadm, kubelet and add a kubelet systemd service:

```
RELEASE="$(curl -sSL
```

```
https://dl.k8s.io/release/stable.txt)" ARCH="amd64"
```

```
cd $DOWNLOAD_DIR
```

```
sudo curl -L --remote-name-all
```

```
https://dl.k8s.io/release/${RELEASE}/bin/linux/${ARCH}/{kubeadm,kubelet
```

```
} sudo chmod +x {kubeadm,kubelet}
```

```
[ec2-user@ip-172-31-41-14 docker]$ RELEASE="$(curl -sSL https://dl.k8s.io/release/stable.txt)"
ARCH="amd64"
cd $DOWNLOAD_DIR
sudo curl -L --remote-name-all https://dl.k8s.io/release/${RELEASE}/bin/linux/${ARCH}/{kubeadm,kubelet}
sudo chmod +x {kubeadm,kubelet}
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	138	100	138	0	0	1643	1662
100	55.5M	100	55.5M	0	0	35.9M	26.1M

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	138	100	138	0	0	3349	3365
100	73.3M	100	73.3M	0	0	89.4M	101M

```
RELEASE_VERSION="v0.16.2"
```

```
curl -sSL
```

```
"https://raw.githubusercontent.com/kubernetes/release/${RELEASE_VERSION}/cmd/kre
```

```
l/templates/latest/kubelet/kubelet.service" | sed "s:/usr/bin:${DOWNLOAD_DIR}:g" |
```

```
sudo tee /usr/lib/systemd/system/kubelet.service
```

```
sudo mkdir -p /usr/lib/systemd/system/kubelet.service.d
```

```
curl -sSL
```

```
"https://raw.githubusercontent.com/kubernetes/release/${RELEASE_VERSION}/cmd/kre
```

```
l/templates/latest/kubeadm/10-kubeadm.conf" | sed
```

```
"s:/usr/bin:${DOWNLOAD_DIR}:g"
```

```
| sudo tee /usr/lib/systemd/system/kubelet.service.d/10-kubeadm.conf
```

```
[ec2-user@ip-172-31-41-14 bin]$ RELEASE_VERSION="v0.16.2"
curl -sSL "https://raw.githubusercontent.com/kubernetes/release/${RELEASE_VERSION}/cmd/krel/templates/latest/kubelet/kubelet.service" | sed "s:/usr/bin:${DOWNLOAD_DIR}:g" | sudo tee /usr/lib/systemd/system/kubelet.service
sudo mkdir -p /usr/lib/systemd/system/kubelet.service.d
curl -sSL "https://raw.githubusercontent.com/kubernetes/release/${RELEASE_VERSION}/cmd/krel/templates/latest/kubeadm/10-kubeadm.conf" | sed "s:/usr/bin:${DOWNLOAD_DIR}:g" | sudo tee /usr/lib/systemd/system/kubelet.service.d/10-kubeadm.conf
[Unit]
Description=kubelet: The Kubernetes Node Agent
Documentation=https://kubernetes.io/docs/
Wants=network-online.target
After=network-online.target

[Service]
ExecStart=/usr/local/bin/kubelet
Restart=always
StartLimitInterval=0
RestartSec=10

[Install]
WantedBy=multi-user.target

# Note: This dropin only works with kubeadm and kubelet v1.11+
[Service]
Environment="KUBELET_KUBECONFIG_ARGS=--bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf"
Environment="KUBELET_CONFIG_ARGS=--config=/var/lib/kubelet/config.yaml"
# This is a file that "kubeadm init" and "kubeadm join" generates at runtime, populating the KUBELET_KUBEADM_ARGS variable dynamically
EnvironmentFile=-/var/lib/kubelet/kubeadm-flags.env
# This is a file that the user can use for overrides of the kubelet args as a last resort. Preferably, the user should use
# the .NodeRegistration.KubeletExtraArgs object in the configuration files instead. KUBELET_EXTRA_ARGS should be sourced from this file.
EnvironmentFile=-/etc/sysconfig/kubelet
ExecStart=
ExecStart=/usr/local/bin/kubelet $KUBELET_KUBECONFIG_ARGS $KUBELET_CONFIG_ARGS $KUBELET_KUBEADM_ARGS $KUBELET_EXTRA_ARGS
[ec2-user@ip-172-31-41-14 bin]$
```

Now we need to install kubect!

Set up repository:

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repo
md.xml.key
EOF
```

```
[ec2-user@ip-172-31-41-14 bin]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
EOF
[ec2-user@ip-172-31-41-14 bin]$
```

```
sudo yum install -y kubectl
```

```
ec2-user@ip-172-31-41-14/us x + v
[ec2-user@ip-172-31-41-14 bin]$ sudo yum install -y kubectl
Kubernetes                                     66 kB/s | 9.4 kB    00:00
Dependencies resolved.
=====
Package                Architecture      Version           Repository        Size
=====
Installing:
  kubectl              x86_64            1.31.1-150500.1.1  kubernetes        11 M
=====
Transaction Summary
=====
Install 1 Package

Total download size: 11 M
Installed size: 54 M
Downloading Packages:
kubectl-1.31.1-150500.1.1.x86_64.rpm          44 MB/s | 11 MB    00:00
-----
Total
Kubernetes                                     42 MB/s | 11 MB    00:00
Importing GPG key 0x9A296436:
  Userid      : "isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
  Fingerprint: DE15 B144 86CD 377B 9E87 6E1A 2346 54DA 9A29 6436
  From        : https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
Key imported successfully
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing      :                                1/1
  Installing     : kubectl-1.31.1-150500.1.1.x86_64 1/1
  Verifying      : kubectl-1.31.1-150500.1.1.x86_64 1/1

Installed:
  kubectl-1.31.1-150500.1.1.x86_64

Complete!
[ec2-user@ip-172-31-41-14 bin]$
```

```
[ec2-user@ip-172-31-41-14 bin]$ kubectl version
Client Version: v1.31.1
Kustomize Version: v5.4.2
```

We have installed successfully installed kubernetes

After installing Kubernetes, we need to configure internet options to allow bridging.

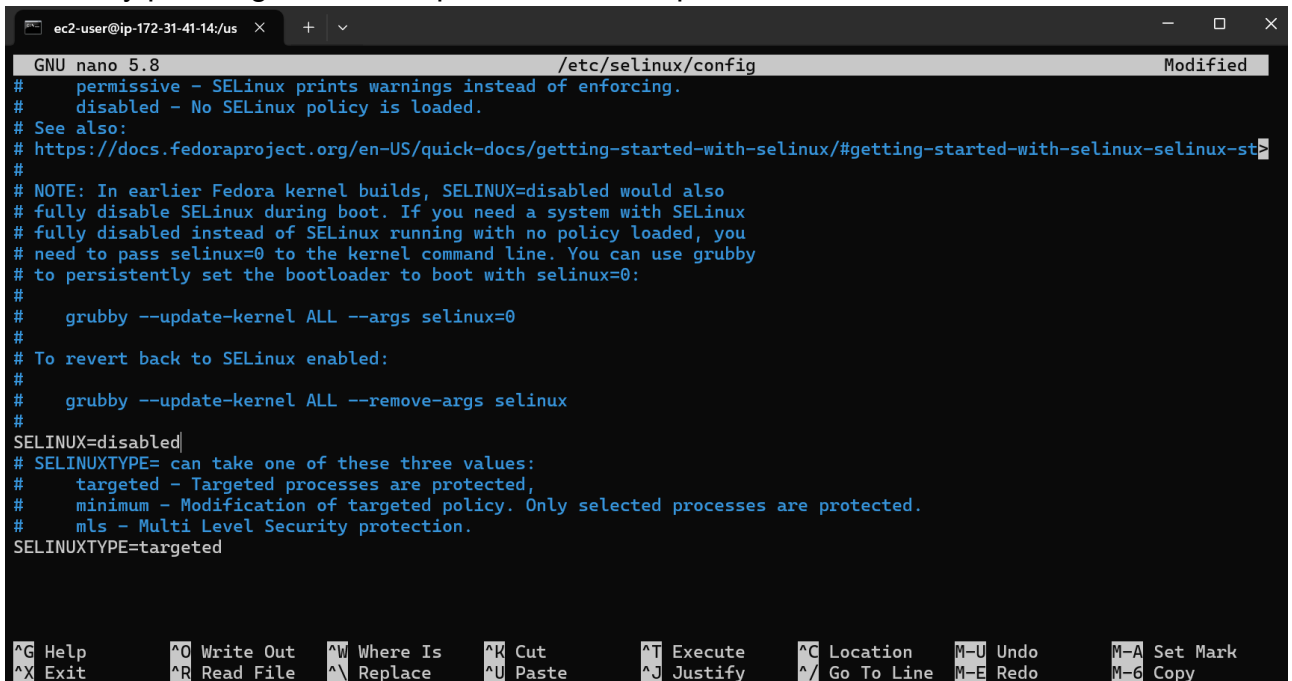
```
sudo swapoff -a
&& echo
"net.bridge.bridge-
e-nf-call-iptables
s=1" | sudo tee
-a
/etc/sysctl.conf
&& sudo sysctl -p
```

```
[ec2-user@ip-172-31-41-14 bin]$ sudo swapoff -a && echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl
.conf && sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-41-14 bin]$
```

Disable SELINUX

Type **sudo nano /etc/selinux/config** and set the value of **SELINUX=disabled** instead of **SELINUX=permissive**

Save the file by pressing ctrl+o then press enter then press ctrl+x



```
ec2-user@ip-172-31-41-14/us  GNU nano 5.8 /etc/selinux/config Modified
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
# See also:
# https://docs.fedoraproject.org/en-US/quick-docs/getting-started-with-selinux/#getting-started-with-selinux-selinux-st
#
# NOTE: In earlier Fedora kernel builds, SELINUX=disabled would also
# fully disable SELinux during boot. If you need a system with SELinux
# fully disabled instead of SELinux running with no policy loaded, you
# need to pass selinux=0 to the kernel command line. You can use grubby
# to persistently set the bootloader to boot with selinux=0:
#
# grubby --update-kernel ALL --args selinux=0
#
# To revert back to SELinux enabled:
#
# grubby --update-kernel ALL --remove-args selinux
#
SELINUX=disabled
# SELINUXTYPE= can take one of these three values:
#   targeted - Targeted processes are protected,
#   minimum - Modification of targeted policy. Only selected processes are protected.
#   mls - Multi Level Security protection.
SELINUXTYPE=targeted

^G Help      ^O Write Out ^W Where Is  ^K Cut       ^T Execute   ^C Location  M-U Undo     M-A Set Mark
^X Exit      ^R Read File ^\ Replace   ^U Paste     ^J Justify   ^/ Go To Line M-E Redo     M-G Copy
```


After rebooting we need to make ssh connection with machine after it gets disconnected

```
[ec2-user@ip-172-31-41-14 bin]$ sudo nano /etc/selinux/config
[ec2-user@ip-172-31-41-14 bin]$ sudo reboot

Broadcast message from root@localhost on pts/1 (Wed 2024-09-18 16:15:38 UTC):

The system will reboot now!

[ec2-user@ip-172-31-41-14 bin]$ Connection to ec2-3-82-9-23.compute-1.amazonaws.com closed by remote host.
Connection to ec2-3-82-9-23.compute-1.amazonaws.com closed.

C:\Users\aditya\Downloads>ssh -i "aditya.pem" ec2-user@ec2-3-82-9-23.compute-1.amazonaws.com

#_
~\_ #####
  \_#####\
  \#####|
  \#/
  \#/_--
  V~'  '--->
      /
     /
    /
   /
  /
 /
/_/m/'

Amazon Linux 2023

https://aws.amazon.com/linux/amazon-linux-2023

Last login: Wed Sep 18 15:59:22 2024 from 202.179.85.199
[ec2-user@ip-172-31-41-14 ~]$
```

Now if we type command **sestatus**, then it show disabled

```
[ec2-user@ip-172-31-41-14 ~]$ sestatus
SELinux status: disabled
[ec2-user@ip-172-31-41-14 ~]$
```

5. Initialize the Kubecluster

Install packages socat and iproute-tc and conntrack to avoid preflight errors

- **sudo dnf install socat iproute-tc conntrack-tools -y**

```

[ec2-user@ip-172-31-41-14:~]$ sudo dnf install socat iproute-tc conntrack-tools -y
Last metadata expiration check: 0:06:41 ago on Wed Sep 18 16:10:55 2024.
Dependencies resolved.
=====
Package                                Architecture      Version            Repository          Size
=====
Installing:
conntrack-tools                        x86_64            1.4.6-2.amzn2023.0.2  amazonlinux        208 k
iproute-tc                             x86_64            5.10.0-2.amzn2023.0.5  amazonlinux        455 k
socat                                  x86_64            1.7.4.2-1.amzn2023.0.2  amazonlinux        303 k
Installing dependencies:
libnetfilter_cthelper                 x86_64            1.0.0-21.amzn2023.0.2  amazonlinux        24 k
libnetfilter_cttimeout                 x86_64            1.0.0-19.amzn2023.0.2  amazonlinux        24 k
libnetfilter_queue                     x86_64            1.0.5-2.amzn2023.0.2  amazonlinux        30 k
=====
Transaction Summary
=====
Install 6 Packages

Total download size: 1.0 M
Installed size: 2.8 M
Downloading Packages:
(1/6): libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64.rpm 461 kB/s | 24 kB 00:00
(2/6): conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64.rpm 3.0 MB/s | 208 kB 00:00
(3/6): libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64.rpm 1.0 MB/s | 24 kB 00:00
(4/6): libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64.rpm 1.2 MB/s | 30 kB 00:00
(5/6): socat-1.7.4.2-1.amzn2023.0.2.x86_64.rpm 5.9 MB/s | 303 kB 00:00
(6/6): iproute-tc-5.10.0-2.amzn2023.0.5.x86_64.rpm 3.1 MB/s | 455 kB 00:00
-----
Total 5.3 MB/s | 1.0 MB 00:00
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing : 1/1
Installing : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 1/6
Installing : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 2/6
Installing : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 3/6
Installing : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 4/6
Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 4/6
Installing : socat-1.7.4.2-1.amzn2023.0.2.x86_64 5/6
Installing : iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 6/6
Running scriptlet: iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 6/6
Verifying : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 1/6
Verifying : iproute-tc-5.10.0-2.amzn2023.0.5.x86_64 2/6
Verifying : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 3/6
Verifying : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 4/6
Verifying : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 5/6
Verifying : socat-1.7.4.2-1.amzn2023.0.2.x86_64 6/6

Installed:
conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 iproute-tc-5.10.0-2.amzn2023.0.5.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 socat-1.7.4.2-1.amzn2023.0.2.x86_64

Complete!

```

- `sudo kubeadm init --pod-network-cidr=10.244.0.0/16`
`--ignore-preflight-errors=NumCPU,Mem`

Copy the mkdir and chown commands from the top and execute them

- `mkdir -p $HOME/.kube`
- `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```
ec2-user@ip-172-31-41-14:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=NumCPU,Mem
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING NumCPU]: the number of available CPUs 1 is less than the required 2
[WARNING Mem]: the system RAM (949 MB) is less than the minimum 1700 MB
[WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl enable kubelet.service'
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0918 16:19:10.818991 3251 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-41-14.ec2.internal kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.41.14]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-41-14.ec2.internal localhost] and IPs [172.31.41.14 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-41-14.ec2.internal localhost] and IPs [172.31.41.14 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
```

- **sudo systemctl restart kubelet**

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.41.14:6443 --token sji44d.lmoyle2o2vxqpty1 \
--discovery-token-ca-cert-hash sha256:e6ea5719242f99612e2a1f595c714c8a123691c05c912e18f6112e90cb67c035
[ec2-user@ip-172-31-41-14 ~]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-41-14 ~]$ sudo systemctl restart kubelet
```

- Then, add a common networking plugin called flannel as mentioned in the code.

kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[ec2-user@ip-172-31-41-14 ~]$
```

Now type **kubectl get nodes**

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-41-14.ec2.internal        Ready     control-plane   3m13s   v1.31.1
[ec2-user@ip-172-31-41-14 ~]$
```

Note: If any time of get error of connection refused just restart the kubelet service (sudo systemctl restart kubelet)

Now that the cluster is up and running, we can deploy our nginx server on this cluster.

Apply this deployment file using this command to create a deployment

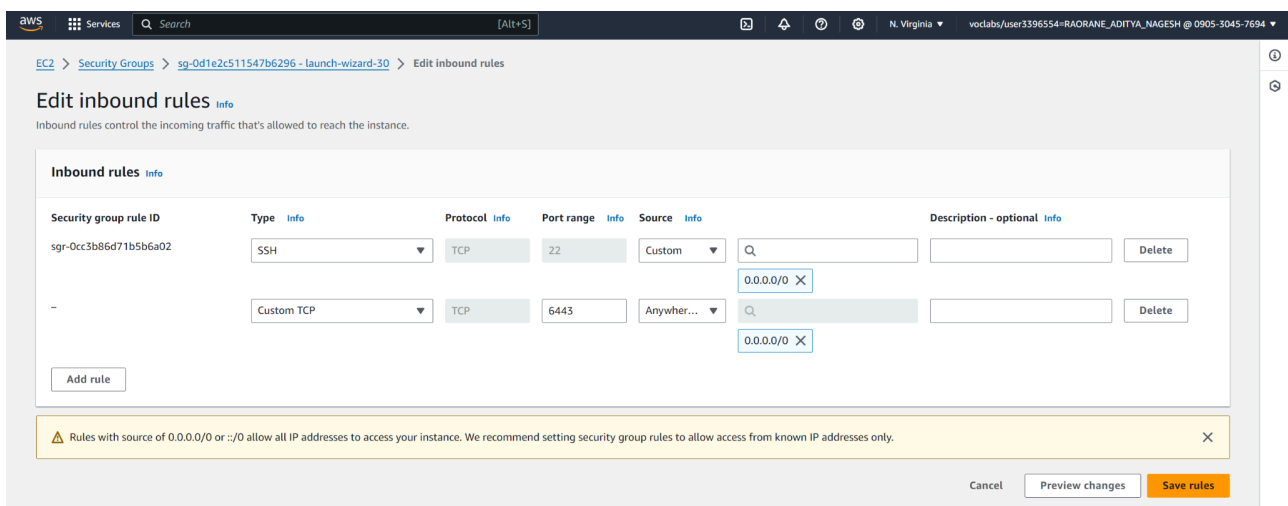
- kubectl apply -f <https://k8s.io/examples/application/deployment.yaml>**

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl apply -f https://k8s.io/examples/application/deployment.yaml
deployment.apps/nginx-deployment created
[ec2-user@ip-172-31-41-14 ~]$
```

Use 'kubectl get pods' to verify if the deployment was properly created and the pod is working correctly.

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl get pods
The connection to the server 172.31.41.14:6443 was refused - did you specify the right host or port?
[ec2-user@ip-172-31-41-14 ~]$
```

Add an inbound rule under SSH security groups which will allow the traffic for a **custom TCP port** with port number **6443** with setting the source as **anywhere from IPv4**.



As we can see our pods are in pending state

On checking logs to we came to know the pods are in tainted state (using command **kubectl describe pod nginx-deployment-d556bf558-7zthh**)

To make pods untainted

Type **kubectl get nodes** to see the name of the node.

Then type command **kubectl taint nodes <NODE_NAME> - -all**

In my case it is as follows:

kubectl taint nodes ip-172-31-41-14.ec2.internal node-role.kubernetes.io/control-plane-

After executing the above command, check again the status of pods if still pending then restart kubelet, wait for 1-2 minutes and check again.

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl taint nodes ip-172-31-41-14.ec2.internal node-role.kubernetes.io/control-plane-  
node/ip-172-31-41-14.ec2.internal untainted
```

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl get pods  
NAME                                READY   STATUS    RESTARTS   AGE  
nginx-deployment-d556bf558-7zthh    1/1     Running   2 (27s ago) 60s  
nginx-deployment-d556bf558-vdldn    1/1     Running   2 (27s ago) 60s  
[ec2-user@ip-172-31-41-14 ~]$ |
```

As we can see our pods are running

- Lastly, port forward the deployment to your localhost so that you can view it. **kubectl port-forward <POD_NAME> 8080:80**

In my case : **kubectl port-forward nginx-deployment-d556bf558-7zthh 8080:80**

Note: if you are getting connection refused error then restart kubelet.

```
[ec2-user@ip-172-31-41-14 ~]$ kubectl port-forward nginx-deployment-d556bf558-7zthh 8080:80  
Forwarding from 127.0.0.1:8080 -> 80  
Forwarding from [::1]:8080 -> 80  
|
```

As port forwarding is active so we cannot type other commands.

Open new terminal window and make ssh connection to same machine

And type command **curl --head <http://127.0.0.1:8080>**

```
[ec2-user@ip-172-31-41-14 ~]$ curl --head http://127.0.0.1:8080
HTTP/1.1 200 OK
Server: nginx/1.14.2
Date: Wed, 18 Sep 2024 16:33:35 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 04 Dec 2018 14:44:49 GMT
Connection: keep-alive
ETag: "5c0692e1-264"
Accept-Ranges: bytes

[ec2-user@ip-172-31-41-14 ~]$ |
```

Response status 200 (OK) indicates that our nginx server is running successfully on kubernetes

Conclusion: We began by installing and configuring Docker and Kubernetes, encountering some initial issues with the Kubernetes API server, which were resolved by restarting the **kubelet** service. The pods didn't start at first due to taints on the nodes, which we removed to allow normal pod scheduling. After resolving these errors, we successfully deployed Nginx server pods and configured them to be accessible via port forwarding. Additionally, we configured the SSH security group by adding an inbound rule to permit traffic on TCP port 6443 from any IPv4 address.