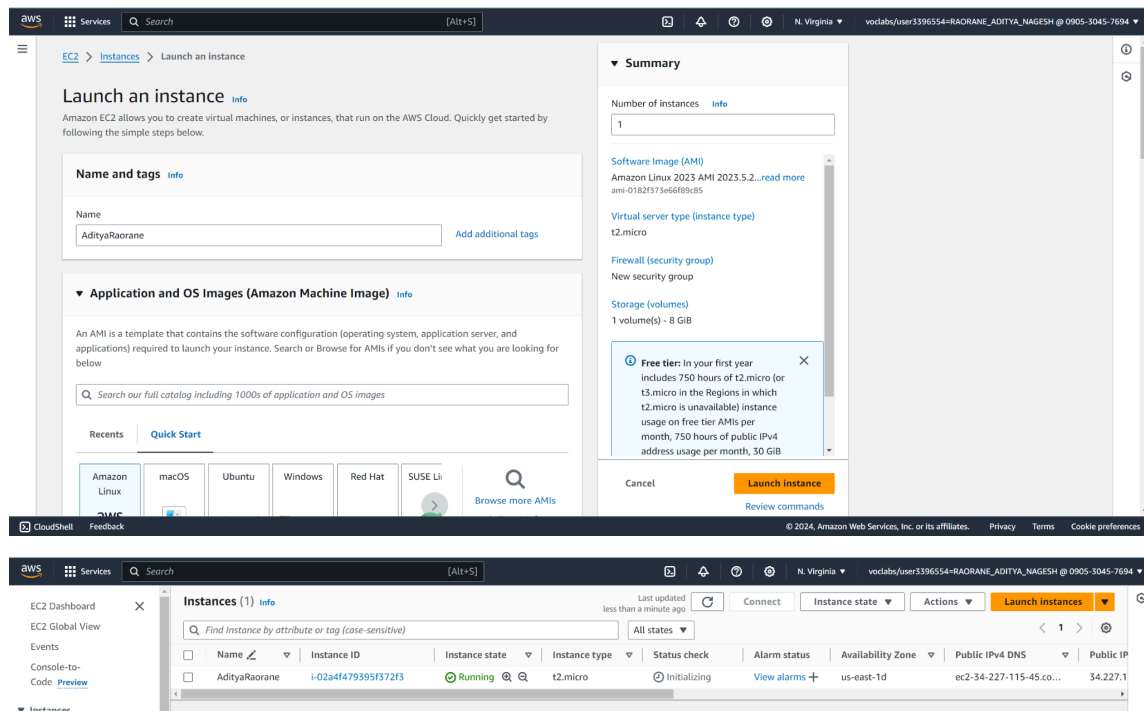


Experiment 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

1. A) Creation Of EC-2 instance →

Launch an AWS EC2 instance named AdityaRaorane using an AWS Linux AMI. Configure the Security Group's Inbound Rules to allow SSH access, then choose the t2.micro instance type.



B| Connecting to an AWS EC2 Instance via SSH →

To connect to an AWS EC2 instance via SSH, change the key file's permission using `chmod 400 "keyname.pem"`, then run `ssh -i <keyname>.pem ubuntu@<public_ip_address>` to establish the connection.

```
ssh -i "master.pem" ec2-user@ec2-34-227-115-45.compute-1.amazonaws.com
```

```
C:\Users\adity\Downloads>ssh -i "master.pem" ec2-user@ec2-34-227-115-45.compute-1.amazonaws.com  
The authenticity of host 'ec2-34-227-115-45.compute-1.amazonaws.com (64:ff9b::22e3:732d)' can't be established.  
ED25519 key fingerprint is SHA256:71B+S9oW46LFCnnIbcyaP5/2VR9+D4Nb8eSlun40r0I.  
This key is not known by any other names  
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes  
Warning: Permanently added 'ec2-34-227-115-45.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
```

```
#  
~# ##### Amazon Linux 2023  
~~#\#####\  
~~~~\###|  
~~~~\#/ https://aws.amazon.com/linux/amazon-linux-2023  
~~~~V--!  
~~~~~_ _ _  
~~~~~_ _ _  
~~~~~/m/_
```

```
[ec2-user@ip-172-31-80-106 ~]$ |
```

2. Installation of Docker →

Run the following command: `sudo yum install docker -y`

```
aws Services Search [Alt+S] N. Virginia vcdlabs/user3396554-RAORANE_ADITYA_NAGESH @ 0905-3045-7694
[ec2-user@ip-172-31-80-106 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:08:10 ago on Sat Sep 14 10:07:59 2024.
Dependencies resolved.
=====
Package                                Architecture      Version            Repository          Size
-----
Installing:
docker                                x86_64            25.0.6-1.amzn2023.0.2  amazonlinux        44 M
Installing dependencies:
containerd                          x86_64            1.7.20-1.amzn2023.0.1  amazonlinux        35 M
iptables-libs                       x86_64            1.8.8-3.amzn2023.0.2  amazonlinux        401 k
iptables-nft                        x86_64            1.8.8-3.amzn2023.0.2  amazonlinux        183 k
libcgroup                           x86_64            3.0-1.amzn2023.0.1    amazonlinux        75 k
libnetfilter_conntrack              x86_64            1.0.8-2.amzn2023.0.2  amazonlinux        58 k
libnftnl                            x86_64            1.0.1-19.amzn2023.0.2  amazonlinux        30 k
libnftnl                            x86_64            1.2.2-2.amzn2023.0.2  amazonlinux        84 k
pigz                                x86_64            2.5-1.amzn2023.0.3    amazonlinux        83 k
runc                                 x86_64            1.1.13-1.amzn2023.0.1  amazonlinux        3.2 M
=====
Transaction Summary
-----
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_64.rpm 6.8 MB/s | 401 kB 00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_64.rpm 4.8 MB/s | 183 kB 00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm 1.1 MB/s | 75 kB 00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64.rpm 1.4 MB/s | 58 kB 00:00
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_64.rpm 823 kB/s | 30 kB 00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rpm 2.8 MB/s | 84 kB 00:00
-----
i-02a4f479395f372f3 (AdityaRaorane)
PublicIPs: 34.227.115.45 PrivateIPs: 172.31.80.106
```

Then, configure cgroup in a daemon.json file by using following commands :

- ❖ `cd /etc/docker`
- ❖ `cat <<EOF | sudo tee /etc/docker/daemon.json`

```
{
  "exec-opts":
  ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
```

```
[ec2-user@ip-172-31-73-36 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

After this run the following command to enable and start docker and also to load the daemon.json file.

- ❖ sudo systemctl enable
- ❖ docker sudo systemctl
- ❖ daemon-reload sudo
- ❖ systemctl restart docker
- ❖ docker -v

```
[ec2-user@ip-172-31-73-36 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-73-36 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-73-36 docker]$ sudo systemctl restart docker
[ec2-user@ip-172-31-73-36 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-73-36 docker]$
```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

3. Install Kubernetes →

a]SELinux needs to be disabled before configuring kubelet Run the following command

- ❖ sudo setenforce 0
- ❖ sudo sed -i 's/^SELINUX=enforcing\$/SELINUX=permissive/' /etc/selinux/config

b] We are adding kubernetes using the repository whose command is given below.

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

```
[ec2-user@ip-172-31-73-36 docker]$ sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-73-36 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[ec2-user@ip-172-31-73-36 docker]$
```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

c] After that Run following command to make the updation and also to install kubelet, kubeadm, kubectl:

- sudo yum update
- sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

```
aws Services Search [Alt+S] N. Virginia voclabs/user3396554=RAORANE_ADITYA_NAGESH @ 0905-3045-7694
[ec2-user@ip-172-31-73-36 docker]$ sudo yum update
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes
Dependencies resolved.
Nothing to do.
Complete!
Last metadata expiration check: 0:00:01 ago on Sat Sep 14 10:49:19 2024.
Dependencies resolved.
```

Package	Architecture	Version	Repository	Size
Installing:				
kubeadm	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubectl	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubelet	x86_64	1.30.5-150500.1.1	kubernetes	17 M
Installing dependencies:				
conntrack-tools	x86_64	1.4.6-2.amzn2023.0.2	amazonlinux	208 k
cri-tools	x86_64	1.30.1-150500.1.1	kubernetes	8.6 M
kubernetes-cni	x86_64	1.4.0-150500.1.1	kubernetes	6.7 M
libnetfilter_cthelper	x86_64	1.0.0-21.amzn2023.0.2	amazonlinux	24 k
libnetfilter_cttimeout	x86_64	1.0.0-19.amzn2023.0.2	amazonlinux	24 k
libnetfilter_queue	x86_64	1.0.5-2.amzn2023.0.2	amazonlinux	30 k
Transaction Summary				
Install 9 Packages				
Total download size: 53 M				
Installed size: 292 M				
Downloading Packages:				
(1/9): libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64.rpm			393 kB/s 24 kB	00:00
(2/9): libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64.rpm			362 kB/s 24 kB	00:00

i-0679633b2451b3c90 (AdityaRaorane)
PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

d] After installing Kubernetes, we need to configure internet options to allow bridging.

1. sudo swapoff -a
2. echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
3. sudo sysctl -p

```
Installed:
conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.30.1-150500.1.1.x86_64      kubeadm-1.30.5-150500.1.1.x86_64
kubectl-1.30.5-150500.1.1.x86_64                kubelet-1.30.5-150500.1.1.x86_64        kubernetes-cni-1.4.0-150500.1.1.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[ec2-user@ip-172-31-73-36 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-73-36 docker]$
```

i-0679633b2451b3c90 (AdityaRaorane)
PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

4. Initialize the Kubecluster →

a] Initializes a Kubernetes cluster with **kubeadm**, sets up the pod network CIDR to **10.244.0.0/16** for network communication, and ignores preflight checks for CPU and memory requirements.

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
--ignore-preflight-errors=NumCPU,Mem
```

```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.73.36:6443 --token xtpltk.qvqcuejymdaszs7b \
  --discovery-token-ca-cert-hash sha256:70773cc7c577bfc8513950c3f619a2c903ff12ad382edc86aa5df024e8c7ecd9
[ec2-user@ip-172-31-73-36 docker]$

```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

b] copy the token and save for future use .

```

kubeadm join 172.31.73.36:6443 --token xtpltk.qvqcuejymdaszs7b \
  --discovery-token-ca-cert-hash
sha256:70773cc7c577bfc8513950c3f619a2c903ff12ad382edc86aa5df024e8
c7ecd9

```

c] Copy the mkdir and chown commands from the top and execute them

- ❖ mkdir -p \$HOME/.kube
- ❖ sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config
- ❖ sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config

```

[ec2-user@ip-172-31-73-36 docker]$ mkdir -p $HOME/.kube
[ec2-user@ip-172-31-73-36 docker]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[ec2-user@ip-172-31-73-36 docker]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

d] Then, add a common networking plugin called flannel as mentioned in the code.

```

kubectl apply -f
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml

```

```

[ec2-user@ip-172-31-73-36 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[ec2-user@ip-172-31-73-36 docker]$

```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

5. Deploying an NGINX Server on Your Kubernetes Cluster →

a] Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply deployment using this following command:

```
kubectl apply -f
```

<https://k8s.io/examples/pods/simple-pod.yaml>

```
[ec2-user@ip-172-31-73-36 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
[ec2-user@ip-172-31-73-36 docker]$
```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

b] Then use **kubectl get pods** to check whether the pod gets created or not.

```
[ec2-user@ip-172-31-73-36 docker]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0           9m33s
[ec2-user@ip-172-31-73-36 docker]$
```

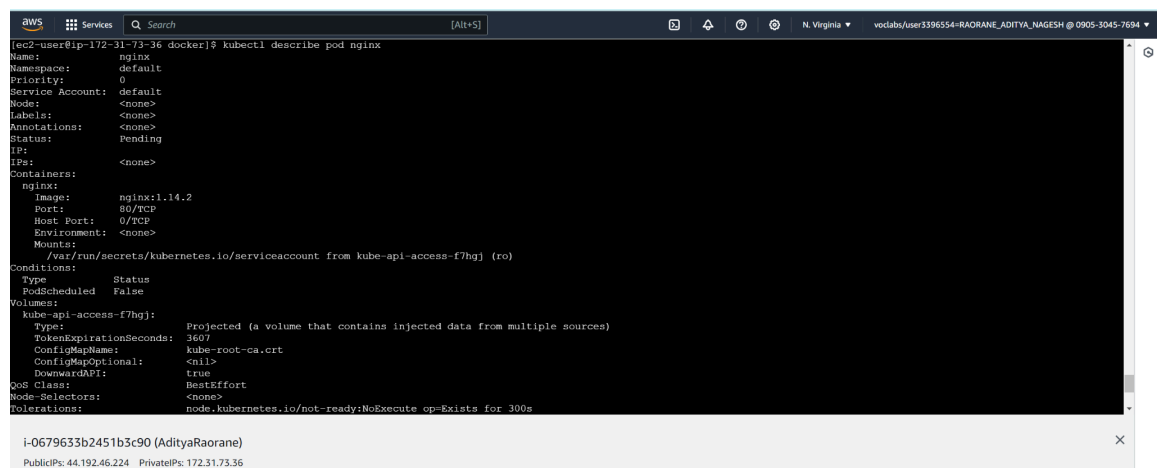
i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

c] To convert state from pending to running use following command:

kubectl describe pod nginx

This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.



```
aws
[ec2-user@ip-172-31-73-36 docker]$ kubectl describe pod nginx
Name: nginx
Namespace: default
Priority: 0
Service Account: default
Node: <none>
Labels: <none>
Annotations: <none>
Status: Pending
IP: <none>
Containers:
  nginx:
    Image: nginx:1.14.2
    Port: 80/TCP
    Host Port: 0/TCP
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-f7hgj (ro)
Conditions:
  Type             Status
  PodScheduled     False
Volumes:
  kube-api-access-f7hgj:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
    BestEffort
  Node-Selectors: <none>
  Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
```

i-0679633b2451b3c90 (AdityaRaorane)

PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

```
[ec2-user@ip-172-31-73-36 docker]$ kubectl taint nodes ip-172-31-73-36.ec2.internal node-role.kubernetes.io/control-plane-  
node/ip-172-31-73-36.ec2.internal untainted
```

- PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

- PublicIPs: 44.192.46.224 PrivateIPs: 172.31.73.36

8. Verify your deployment

Open up a new terminal and ssh to your EC2 instance.

Then, use this curl command to check if the Nginx server is running.

```
curl --head http://127.0.0.1:8080
```

If the response is 200 OK and you can see the Nginx server name, your deployment was successful. We have successfully deployed our Nginx server on our EC2 instance.

Conclusion:

Thus we established a Kubernetes cluster on AWS EC2, configured Docker with cgroup settings, and set up Kubernetes components. We applied Flannel for networking and deployed an NGINX server. Next, we exposed the NGINX server using a Kubernetes service to allow external access. Finally, we implemented autoscaling based on CPU usage to dynamically manage traffic load on the NGINX pods.